

GTEx multi-tissue eQTL analysis

Gao Wang

Last updated: November 13, 2015

Chapter 1

General Information

1.1 Release & Analysis V4

Obsolete

- GTEx raw and normalized expression and PEER factors data:
 - ▶ ftp://ftp.broadinstitute.org/GTEx_Analysis_2014-06-13/expressionCovariateSandbox/
 - ▶ There is a README file in this folder that contains description of data

- GTEx summary statistics formatted to eqtlbma input, by Sarah:

Midway </project/mstephens/gtex/preprocessing/june2014/sumstats>

- GTEx summary statistics from eqtlbma by Sarah:

Midway </project/mstephens/gtex/analysis/june2014/sumstats>

- GTEx genotypes
 - ▶ </project/mstephens/gtex/preprocessing/june2014/omni451.vcf.gz>
- GTEx expression level

Midway /project/mstephens/gtex/preprocessing/june2014/GTEx_Analysis_2014-06-13_expression_provisional/

1.2 Release V6

Input and summary statistics from MatrixEQTL. See [this page](#)¹ for details.

¹this page <https://github.com/stephenslab/lab-resource/blob/master/data/gtex-v6-eqtl.md>

1.2.1 Why covariates not removed from expression values

“One reason for this is because this is how Matrix eQTL accepts input data. I realize that corrected expression values could be valuable to people and would certainly be willing to provide them, but I do not have a method. Although PEER does output a residuals file, I can only get PEER to run robustly by filtering low expressed genes and opting not to use explicit covariates as input to PEER. The solution would be to use the transformation that Matrix eQTL uses internally. ...” [David DeLuca, Broad institute]

Chapter 2

GTEx Summary Statistics

Computations are done on the Midway cluster. `GTEx.bashrc` sets up the proper shell environment for this computational workflow:

```
source $HOME/GIT/gtex-eqtl/src/cfg/GTEx.bashrc
```

2.1 Formatting GTEx Summary Statistics V6

2.1.1 Convert tissue specific results to matrices

For ease of storage / access / numeric operations I convert the summary statistics from plain text file to HDF5 format.

ss.gz to HDF5 conversion

```
cd $AnlyDir/GTEx_Analysis_2015-01-12_MatrixEQTL_allCisSNPGenePairs
for i in `ls *.gz`; do
echo `#!/bin/bash
    source $HOME/GIT/gtex-eqtl/src/cfg/GTEx.bashrc
    python $SrcDir/python/analysis_admin.py ss_to_h5 \
    "$i" --action convert --output . --message "MatrixEQTL summary statistics of GTEx Release 2015.04.15" | \
    sbatch -J ss2h5_${i} -o $LogDir/ss2h5_${i}.o%j -e $LogDir/ss2h5_${i}.e%j
done`
```



Note

I ended up having to perform the entire data conversion on SSD on my desktop computer using. The pytable version of HDF5 implementation takes 2h to convert each tissue but on Midway when HDF5 file size grows to over 1GB the disk I/O is unacceptable. There must be tweaks but not looking into it for the moment. Implementations with h5py for the same purpose is way faster in data conversion (30min each file, performance on Midway is also acceptable) but the resulting data file size is 1.4 times of pytables implementation, even both set to zlib9 compression – thus I stick with pytable. Once converted, it takes almost zero seconds to query data from the HDF5 files.

2.1.2 Merge data across tissues

HDF5 merge, in 50 batches

```
source $HOME/GIT/gtex-eqtl/src/cfg/GTEEx.bashrc
zcat $InputDir/tss_coords.bed.gz | cut -f4 | split -l 1127 -d - $TmpDir/tss_coords.bed.split.
k=0
for i in `ls $TmpDir/tss_coords.bed.split.*`; do
let k+=1
echo `#!/bin/bash
source $HOME/GIT/gtex-eqtl/src/cfg/GTEEx.bashrc
python $SrcDir/python/analysis_admin.py ss_to_h5 \
$DataDir/GTEEx_Analysis_2015-01-12_MatrixEQTL_allCisSNPGenePairs/*.h5 \
--action merge --output $DataDir/MatrixEQTLSumStats/MatrixEQTLSumStats."$k" \
--message "MatrixEQTL summary statistics of GTEx Release 2015.04.15" \
--gene-list "$i" | \
sbatch -J h5merge_"$k" -o $LogDir/h5merge_"$k".o%j -e $LogDir/h5merge_"$k".e%j
done`
```

It takes on average 3 hours per batch to complete. Merged number of genes per batch see log/2015-06-18-merger-log.txt. There are a total of 38933 genes.



Note

I was suspicious to see only 38933 genes in v6. I decide to check against the input data as well as the intermediate HDF5 files I generated

BASH

```
for i in `ls $DataDir/GTEEx_Analysis_2015-01-12_MatrixEQTL_allCisSNPGenePairs/*.cis.eqtl.gz`; do
echo `#!/bin/bash
source $HOME/GIT/gtex-eqtl/src/cfg/GTEEx.bashrc
zcat "$i" | cut -f2 | sort -u | gzip --best > "$i".uniqgenes.gz | \
sbatch -J count_gene -o $LogDir/count_gene.o%j -e $LogDir/count_gene.e%j
done
for i in `ls $DataDir/GTEEx_Analysis_2015-01-12_MatrixEQTL_allCisSNPGenePairs/*.h5`; do
echo `#!/bin/bash
source $HOME/GIT/gtex-eqtl/src/cfg/GTEEx.bashrc
python $SrcDir/python/analysis_admin.py ss_to_h5 "$i" --action summary --output stdout | gzip --best > "$i".h5genes.gz | \
sbatch -J count_gene_h5 -o $LogDir/count_gene_h5.o%j -e $LogDir/count_gene_h5.e%j
done`
```

Counting (zcat *.xx.gz | sort -u | wc -l) and comparing results from both sources, I found there are indeed only 38933 genes in the output. So we are good.

Now merging the batches into one large HDF5 file

```
python $SrcDir/python/analysis_admin.py ss_to_h5 $DataDir/MatrixEQTLSumStats/*.h5 --action cat --output $DataDir/MatrixEQTLSumStats.h5
```

2.1.3 Select the “best” gene-snp pair

```
python $SrcDir/python/analysis_admin.py ss_to_h5 $DataDir/MatrixEQTLSumStats.h5 --action max --output $DataDir/MatrixEQTLSumStats.Portable.h5
```

2.1.4 Select some “null” gene-snp pairs

For each gene I select up to 3 null gene-snp pairs, for use of calibrating Sarah’s model

```
python $SrcDir/python/analysis_admin.py ss_to_h5 $DataDir/MatrixEQTLSumStats.h5 --action null --output $DataDir/Mat\
rixEQTLSumStats.Portable.h5 --gene-list <(zcat $DBDir/snp-gene-pairs/GTEXV6Genes.gz)
```

2.2 Creating R Interface for the New Format

Since the release did not use rsID to name SNPs I create a database to match SNP ID with dbSNP names, to use with R to extract proper SNP-gene by rsID.

```
wget ftp://ftp.ncbi.nih.gov/snp/organisms/human_9606_b144_GRCh37p13/VCF/All_20150605.vcf.gz
```

Build an annotation database matching SNP ID to rs ID as well as the cis-genes involved of up to $\pm 100,000$ bp.

ID annotations, in 50 batches

```
source $HOME/GIT/gtex-eqtl/src/cfg/GTEX.bashrc
zcat $InputDir/snp_coords.bed.gz | split -l 205953 -d - $TmpDir/snp_coords.bed.split.
k=0
for i in `ls $TmpDir/snp_coords.bed.split.*`; do
let k+=1
echo `#!/bin/bash
source $HOME/GIT/gtex-eqtl/src/cfg/GTEX.bashrc
cat `"$i"` | bash $SrcDir/shell/FindrsIDcisGene.sh \
$DBDir/dbSNP/All_20150605.vcf.gz $InputDir/tss_coords.bed.gz 100000 | gzip > \
$DBDir/snp-gene-pairs.`"$k"`.gz' | \
sbatch -J snpid_`"$k"` -o $LogDir/snpid_`"$k"`.o%j -e $LogDir/snpid_`"$k"`.e%j
done
```

It takes on average 25 hours per batch to complete. There are 9794339 out of 10297646 (95.1%) SNPs with rsID in latest dbSNP (build 144). To merge them to one file and build a sqlite database:

```
zcat $DBDir/snp-gene-pairs.*.gz | gzip --best > $DBDir/snp-gene-pairs.gz
cd $DBDir; mkdir snp-gene-pairs; mv snp-gene-pairs.* snp-gene-pairs; cd snp-gene-pairs
sqlite snp-gene.sqlite3 -i snp-gene-pairs.gz --as dbsnp144 -d '\t' --header coord rsid cisgenes
sqlite snp-gene.sqlite3 "create index rsid_index on dbsnp144 (rsid)"
sqlite snp-gene.sqlite3 "create index coord_index on dbsnp144 (coord)"
```

2.2.1 Run queries

Now we have the data *.h5 and meta-data *.sqlite in place. I have made a separate note on how to use this data-set. See SumstatsDB.pdf in this repo.

2.3 Reproducing GTEx V6 Cis-eQTL Analysis

The idea now is to figure out how the matrixEQTL analysis is done exactly in v6 so that eqtlbma will be done on the same basis, i.e., input files and parameters. For this purpose I make a toy dataset for a randomly chosen tissue from GTEx consisting of five files for input to matrixEQTL:

genotype SNP.txt, expression GE.txt, a file Covariates.txt of covariates and files geneloc.txt and snpsloc.txt with gene and SNP location information.

```
Tissue=Adipose_Subcutaneous
```

2.3.1 Covariates

There are 40 covariates: “C1,C2,C3,InferredCov1, InferredCov2, ... InferredCov35, gender, Platform”. Not all are included in every analysis – inclusion depends on sample size. See the README for details.

```
tar -zxvf GTEx_Analysis_2015-01-12_eQTLInputFiles_covariates.tar.gz "$Tissue"_Analysis.covariates.txt
wc -l "$Tissue"_Analysis.covariates.txt
cut -f1 "$Tissue"_Analysis.covariates.txt | tr "\n" " "
mv "$Tissue"_Analysis.covariates.txt $TmpDir/Covariates.txt
```

2.3.2 SNP data

■ SNP matrix

Input genotype data are large. Here I just take the first 500 SNPs from the tissue data for verification purpose.

```
tar -zxvf GTEx_Analysis_2015-01-12_eQTLInputFiles_snpMatrices.tar.gz "$Tissue"_Analysis.snps.txt -0 |\
head -501 > $TmpDir/SNP.txt
```



Warning

SNP data released has a lot of imputed data. Cautions are to be taken when comparing with previous eqtlbma analysis (have to make sure the input data are the same).

■ SNP location

This is required input for matrixEQTL but is not provided. Need to create it from SNP ID's.

```
echo -e "snp\tchr\tpos" > $TmpDir/snpsloc.txt
cut $TmpDir/SNP.txt -f1 | tail -n+2 | awk -F"_" '{print $0"\t"chr"$1"\t"$2}' >> $TmpDir/snpsloc.txt
```

2.3.3 Expression data

To focus only on genes paired with the 500 SNPs I get a list of genes via bedtools closest and only extract expression data for these genes.

Create bedtools input and run bedtools to search for genes within 1000bp range to the SNPs of interest.

```
tail -n+2 GTEx_Analysis_2015-01-12_eQTLInputFiles_genePositions.txt | \
awk '{print $2"\t"$3"\t"$4+1"\t"$1"\t"1"\t"+">' > $TmpDir/geneloc.bed
cut $TmpDir/SNP.txt -f1 | tail -n+2 | awk -F"_" '{print $1"\t"$2"\t"$2+1"\t"$0"\t"1"\t"+">' > $TmpDir/snpsloc.bed
bedtools closest -a $TmpDir/snpsloc.bed -b $TmpDir/geneloc.bed | \
awk '{if (($8-$2)<1000 && ($8-$2)>-1000) print $0}' > $TmpDir/gs-pairs.bed
```

Extract expression data for selected genes

```
cut -f10 $TmpDir/gs-pairs.bed | sort -u > $TmpDir/1.txt # 22 genes found
tar -zxvf GTEx_Analysis_2015-01-12_eQTLInputFiles_geneLevelNormalizedExpressionMatrices.tar.gz \
"$Tissue"_Analysis.expr.txt -O | head -1 > $TmpDir/GE.txt
tar -zxvf GTEx_Analysis_2015-01-12_eQTLInputFiles_geneLevelNormalizedExpressionMatrices.tar.gz \
"$Tissue"_Analysis.expr.txt -O | grep -wf $TmpDir/1.txt >> $TmpDir/GE.txt # 11 genes here
```

It ends up with 11 genes for the 500 SNPs

```
cut -f 1,2 $TmpDir/GE.txt
```

2.3.4 Gene location data

Location data is extracted for the selected genes.

```
cut -f1 $TmpDir/GE.txt | tail -n+2 > $TmpDir/1.txt
grep -wf $TmpDir/1.txt GTEx_Analysis_2015-01-12_eQTLInputFiles_genePositions.txt > $TmpDir/2.txt
head -1 GTEx_Analysis_2015-01-12_eQTLInputFiles_genePositions.txt | cat - $TmpDir/2.txt > $TmpDir/geneloc.txt
```

2.3.5 Find overlapping gene-snp pairs between this toy set and GTEx v6 summary statistics

BASH

```
for i in `cut -f1 $TmpDir/geneloc.txt | tail -n+2`; do
  echo $i
  python analysis_admin.py sumstat_query $InputDir/"$Tissue"_Analysis.h5 -g $i -s NULL > $TmpDir/1.txt
  if [[ -s $TmpDir/1.txt ]]; then
    grep -wf $TmpDir/1.txt $TmpDir/snpsloc.txt | cut -f1 > $TmpDir/"$i"_snps_to_check.txt
  fi
done
```

2.3.6 Run matrixEQTL on this toy set

Once the input is prepared it is straightforward to just download the example R script (below, also committed to git repo) and run it, just remember to set the correct data path and change emitted p-value threshold to 1 (pvOutputThreshold=1).

http://www.bios.unc.edu/research/genomic_software/Matrix_eQTL/R.html

To obtain value for one gene-snp pair:

R

```
source('MatrixEQTL.R')
ss = '1_714019_A_G_b37'
gg = 'ENSG00000237491.4'
me$all$eqtls[which(me$all$eqtls$snps == ss & me$all$eqtls$gene == gg ),]
```

	snps	gene	statistic	pvalue	FDR
241	1_714019_A_G_b37	ENSG000000237491.4	5.392883	1.576441e-07	3.597687e-06
	beta				
241	0.7902205				

and to verify value for gene-snp pairs:

```
echo $Tissue
```

```
source('/project/mstephens/gtex/scripts/SumstatQuery.R')
dat <- GetFlatSS('ENSG00000237491.4', '/project/mstephens/gtex/analysis/april2015/eqt1_data/GTex_Analysis_2015-01-12_MatrixEQTL_allC
isSNPGenePairs/Adipose_Subcutaneous_Analysis.h5')
print(dat["1_714019_A_G_b37", ])
```

	beta	t-stat	p-value
7.902205e-01	5.392883e+00	1.576441e-07	

I verified a few other gene-SNP pairs, all agree. So I can claim I've figured out the input data and parameters used for the GTEx analysis!

2.4 Update eqtlbma Analysis for Some Gene-snp Pairs and Verify with GTEx matrixEQTL

Here is the command for generating summary statistics via eqtlbma_bf:

```
Analysis=join
Tissue=Test_Tissue
echo -e "$Tissue\tSNP.txt.gz" > list_genos.txt
echo -e "$Tissue\tGE.txt" > list_expl.txt
echo -e "$Tissue\tCovariates.txt" > list_covar.txt
eqtlbma_bf --geno list_genos.txt --exp list_expl.txt --covar list_covar.txt \
--scoord snpsloc.bed --gcoord geneloc.bed \
--out ToyEQTLBma --analys $Analysis --outss --bfs sin \
--gridL /project/mstephens/gtex/analysis/june2014/types_v1/grid_phi2_oma2_general.txt.gz \
--gridS /project/mstephens/gtex/analysis/june2014/types_v1/grid_phi2_oma2_with-configs.txt.gz \
--error uvlr --thread 4
```

Check a couple of output from eqtlbma_bf with matrixEQTL

```
zcat ToyEQTLBma_sumstats_Test_Tissue.txt.gz | head | cut -f1,2,7,9
```

gene	snp	betahat.geno	betapval.geno	
ENSG000000177757.1		1_662622_G_A_b37	-2.588170e-04	9.985439e-01
ENSG000000177757.1		1_676127_C_T_b37	7.151412e-02	7.744839e-01
ENSG000000177757.1		1_691541_AT_A_b37	-9.077870e-02	4.383949e-01

ENSG00000177757.1	1_693625_T_C_b37	-5.757815e-02	7.415218e-01
ENSG00000177757.1	1_693731_A_G_b37	-1.080375e-02	9.255429e-01
ENSG00000177757.1	1_697411_G_GA_b37	5.502301e-01	2.327584e-02
ENSG00000177757.1	1_701835_T_C_b37	7.243649e-01	3.168885e-03
ENSG00000177757.1	1_704367_T_C_b37	1.898470e-01	3.756234e-01
ENSG00000177757.1	1_705882_G_A_b37	-1.595642e-01	3.141656e-01

Check with MatrixEQTL

R

```
gg = 'ENSG00000177757.1'
for (ss in c('1_662622_G_A_b37', '1_676127_C_T_b37', '1_691541_AT_A_b37', '1_693625_T_C_b37', '1_693731_A_G_b37')) print(me$all$eqtl
s[which(me$all$eqtls$snps == ss & me$all$eqtls$gene == gg ),J[,c(2,1,6,4)])
```

OUTPUT

	gene	snps	beta	pvalue
5493	ENSG00000177757.1	1_662622_G_A_b37	-0.000258817	0.9985439
	gene	snps	beta	pvalue
4593	ENSG00000177757.1	1_676127_C_T_b37	0.07151412	0.7744839
	gene	snps	beta	pvalue
3047	ENSG00000177757.1	1_691541_AT_A_b37	-0.0907787	0.4383949
	gene	snps	beta	pvalue
4454	ENSG00000177757.1	1_693625_T_C_b37	-0.05757815	0.7415218
	gene	snps	beta	pvalue
5200	ENSG00000177757.1	1_693731_A_G_b37	-0.01080375	0.9255429

Thus eqtlbma.bf agrees with MatrixEQTL when input agree. Good! We'll have to update the analysis using new input data.

GTEx Summary Statistics Database

3.1 Databases

3.1.1 Multi-tissue summary stats

Summary statistics from the GTEx v6 analysis are stored in `MatrixEQTLSumStats.h5` (130GB, in HDF5 format). It contains 38,933 genes (less than previous version! See the list `/project/mstephens/data`) each gene has several summary statistics matrix of dimension $N_{cisSNPs} \times N_{tissues}$. For this release $N_{tissues} = 44$. For `MatrixEQTL` the summary statistics are $\hat{\beta}$, T -stat and P -value.

`MatrixEQTLSumStats.h5`

```
/ENSG000000008735.10 'MatrixEQTL summary statistics of GTEx Release 2015.04.15'
/ENSG000000008735.10/beta (5636, 44)
/ENSG000000008735.10/p-value (5636, 44)
/ENSG000000008735.10/t-stat (5636, 44)
/ENSG000000008735.10/rownames (5636,)
/ENSG000000008735.10/colnames (44,)
...
```

There are two special tables in `MatrixEQTLSumStats.h5`. One called `max`, which is the data for the “best” gene-snp pair identified per gene as the snp pair having maximum `abs(t-stat)` across all snps and all tissues and there is no missing data for any tissue. It contains 16,069 gene-snp pairs. The other called `null`, which organize gene-snp pair data in the same format as `max` only that the gene-snp pairs in this table are NOT the best gene-snp pair, but just random samples from all non-best gene-snp pairs. For portability these two tables are also stored separately in `MatrixEQTLSumStats.Portable.h5` (58MB, in HDF5 format).

`MatrixEQTLSumStats.Portable.h5`

```
/max (Group) 'max'
/max/beta (16069, 44)
/max/p-value (16069, 44)
/max/t-stat (16069, 44)
/max/rownames (16069,)
/max/colnames (44,)
```

3.1.2 Meta-database

Unfortunately GTEx names *cis*SNP by their genomic coordinate in Human Genome Assembly hg19 (b37), not by rsID. To make it possible to search with rsID I created another database `snp-gene.db` (1.6GB, in SQLite format).

snp-gene.db		
3_52771872_C_T_b37	rs59638016	ENSG00000016864.12,ENSG00000055955.11,ENSG00000055957.6 ...
3_52774530_C_T_b37	rs2159644	ENSG00000016864.12,ENSG00000055955.11,ENSG00000055957.6 ...
3_52774835_AT_A_b37	rs35270761,rs397785457	ENSG00000016864.12,ENSG00000055955.11,ENSG00000055957.6 ...
...		

The 3 columns are

- SNP name in GTEx convention
- rsID
- cis-genes, *i.e.*, genes whose TSS is within 100,000bp up/down-stream of the SNP.

This database will be used prior to searching the summary statistics.



Note

- There may be multiple rsID associated with the same coordinate. Indeed such records exists in dbSNP 144, and eventually these rsIDs should merge into one name (as always have happened in history).
- There are 10,297,646 SNPs in total in this dataset. 9,794,339 of them (95.1%) has rsID in dbSNP 144.

3.2 Queries

I provide a simple R script `SumstatQuery.R` to make queries on this dataset. It has:

- A function that lists for given rsID the GTEx SNP ID and all its cis-genes.
- A function that loads sumstats data given gene name and GTEx SNP ID.

To use the script, `RSQLite` and `rhdf5` should be installed, *e.g.*,

```
install.packages("RSQLite")
source("http://bioconductor.org/biocLite.R")
biocLite("rhdf5")
```

3.2.1 Example query

The demo script below shows an example to extract information for gene ENSG00000171960.6 and SNP rs6600419, and examples to extract the best/null gene-snp pairs.

R

```
source("/project/mstephens/gtex/scripts/SumstatQuery.R")
##
# Load data for given gene
##
dat <- GetSS("ENSG00000171960.6", "/project/mstephens/gtex/analysis/april2015/query/MatrixEQTLSumStats.h5")
print(names(dat))
## [1] "beta" "p-value" "t-stat" "z-score"
dim(dat$"beta")
## [1] 6344 44
dat$"beta"[1:4,1:4]
##
## Adipose_Subcutaneous Adipose_Visceral_Omentum Adrenal_Gland Artery_Aorta
## 1_42124161_C_T_b37 0.04608095 NaN NaN 0.06785597
## 1_42124246_A_AC_b37 -0.03087348 -0.03222961 -0.1286249 -0.13490575
## 1_42124311_G_A_b37 0.12400229 0.08011222 NaN -0.21858459
## 1_42124614_C_T_b37 0.04630340 NaN NaN 0.06785597
##
# Output information for the SNP of interest
##
dat$"p-value"["1_43124701_A_G_b37", ]
##
## Adipose_Subcutaneous Adipose_Visceral_Omentum
## 1.632499e-01 3.462833e-02
## Adrenal_Gland Artery_Aorta
## 8.168959e-01 5.600338e-01
## Artery_Coronary Artery_Tibial
## 7.659426e-01 7.397466e-01
## Brain_Anterior_cingulate_cortex_BA24 Brain_Caudate_basal_ganglia
## ...
dat$"z-score"["1_43124701_A_G_b37", ]
##
## Adipose_Subcutaneous Adipose_Visceral_Omentum
## -1.39422418 -2.11267803
## Adrenal_Gland Artery_Aorta
## -0.23153929 0.58279135
## Artery_Coronary Artery_Tibial
## 0.29768626 0.33218889
## Brain_Anterior_cingulate_cortex_BA24 Brain_Caudate_basal_ganglia
## ...
###
# Load the best gene-snp data
###
mdat <- GetSS("max", "/project/mstephens/gtex/analysis/april2015/query/MatrixEQTLSumStats.h5")
dim(mdat$"t-stat")
mdat$"p-value"[1:4,1:4]
mdat$"t-stat"["ENSG00000000419.8_20_49461813_G_C_b37",]
## Is this gene-snp pair most significant in spleen?
dat <- GetSS("ENSG00000000419.8", "/project/mstephens/gtex/analysis/april2015/query/MatrixEQTLSumStats.h5")
idx.to.show <- matxMax(abs(dat$"t-stat"))
rownames(dat$"t-stat")[idx.to.show[1]]
colnames(dat$"t-stat")[idx.to.show[2]]
###
# Load the "null" gene-snp data
###
ndat <- GetSS("null", "/project/mstephens/gtex/analysis/april2015/query/MatrixEQTLSumStats.h5")
dim(ndat$"t-stat")
##
# Look up GTEx SNP ID
##
ShowSNP("rs6600419", "/project/mstephens/gtex/analysis/april2015/query/snp-gene.db")
## GTEx SNP ID: 1_43124701_A_G_b37
## cisGenes: ENSG000000065978.13,ENSG00000164007.6,ENSG00000171960.6,ENSG00000200254.1,ENSG00000234917.1,ENSG00000236180.2
##
# Look up rs ID given GTEx SNP ID
##
ShowSNP("1_43124701_A_G_b37", "/project/mstephens/gtex/analysis/april2015/query/snp-gene.db")
## rsID(s): rs6600419
## cisGenes: ENSG000000065978.13,ENSG00000164007.6,ENSG00000171960.6,ENSG00000200254.1,ENSG00000234917.1,ENSG00000236180.2
##
# Create matched training/testing sets
##
```

```

N1 <- 8000
N2 <- 16069
strong.train <- SubsetMatLists(mdat, seq(1, N1))
strong.test <- SubsetMatLists(mdat, seq(N1 + 1, N2))
strong.train.genes <- as.character(lapply(strsplit(rownames(strong.train$beta), "_"), function(x) x[1]))
strong.test.genes <- as.character(lapply(strsplit(rownames(strong.test$beta), "_"), function(x) x[1]))
null.genes <- as.character(lapply(strsplit(rownames(ndat$beta), "_"), function(x) x[1]))
null.train <- SubsetMatLists(ndat, which(null.genes %in% strong.train.genes))
null.test <- SubsetMatLists(ndat, which(null.genes %in% strong.test.genes))

```

For `MatrixEQTLSumStats.h5`, for given `rsID` you first search the GTEx SNP ID via `ShowSNP`, then load the `cisGene` of interest via `GetSS` and extract information using the GTEx SNP ID as the row name key.



Note

If R complains the given row name does not exist, try to flip the DNA strand in the GTEx SNP ID and search again. For example, instead of searching for `1_43124701_A_G.b37` you search for `1_43124701_T_C.b37`.

The best gene-snp data can be loaded to memory, via `GetSS('max', 'MatrixEQTLSumStats.h5')`, or the portable version `GetSS('max', 'MatrixEQTLSumStats.Portable.h5')`

Same for the “null” gene-snp data, via `GetSS('null', 'MatrixEQTLSumStats.h5')`, or the portable version `GetSS('null', 'MatrixEQTLSumStats.Portable.h5')`

Computational Steps for Calculating Summary Statistics Under Quasi-Poisson Model

4.1 First Attempt

Here is the command which breaks the data into 2,000 batches and perform analysis under quasi-Poisson model.

2000 batches analysis

```
nBatches=2000
Model=quasipoisson
for i in `seq $nBatches`; do
echo `#!/bin/bash
source $HOME/GIT/type-model/conf/GTEx.bashrc
python $SrcDir/analysis_admin.py eqtlbma_batch \
-g $InputDir/tss_coords.bed.gz \
-s $InputDir/snp_coords.bed.gz \
-n "$nBatches" -b "$i" -e ~/software/bin/eqtlbma_bf \
-a $ConfDir/eqtlbma."$Model".txt' | \
sbatch -J eqtlbma_`Model`$nBatches_$i -o $LogDir/eqtlbma_`Model`$nBatches_$i.o%j --mem-per-cpu=10000
done
```

4.1.1 Result

Computation is heavy and not completed (see next section). Here taking a sample batch out of the 2000 batches above, the following is summarized from eqtlbma log file

- Input batch has 28 genes and 113 SNPs
- Genotype and expression of 44 tissues are loaded, vary in sample size, from >70 to 450.
- Tissue specific genes “to keep” vary, but are around 26, 27, 28.
- Number of SNPs (from omni451.vcf.gz) drop greatly after discarding missing values and filter by $MAF < 0.05$. There are only 10 out of 113 SNPs left.

- Test for each gene-snp pair are done for every batch. There are only a handful such pairs for 2000 batches.

When 5000 batches are created in another test run, some batch may end up fail after loading the data because there will be no qualified SNP left for that batch, from our sample data.

4.1.2 A note on computational time

There will be a 20min “overhead” for each job (loading / extracting data) regardless of per batch size, i.e., we will “waste” 0.3N CPU hours for running in N batches. It is thus best to find a batch neither too small nor too large (if we do not optimize the way data is loaded).

In my test job when only 10 SNPs and <28 genes are analyzed in effect, it takes only 1.38 sec to complete for normal model, evaluating for each tissue about 200 gene-SNP pairs.

However for the same job under quasi-Poisson it takes **significantly** computation time. For example 5 hours (10:30 am to 3:30 pm) have passed for the 10 SNPs and 28 genes. Still the progress bar is at 3.75%. I have tried --thread option of eqtlbma_bf and reserve corresponding number of nodes per job on Midway. Somehow the computational node still only use one thread for a job. May have to figure it out but I doubt it's going to help now that we can easily create arbitrary batches.

Here is the eqtlbma_bf command in action for this one batch. For this batch genes_2000_85-112.bed.gz has 28 lines and snps_2000_85_112.txt.gz has 113 lines.

Quasi-Poisson eqtlbma_bf command

```
/home/gaow/software/bin/eqtlbma_bf --geno /project/mstephens/gtex/analysis/june2014/types_v1/list_geno.txt \
--exp /project/mstephens/gtex/analysis/june2014/types_v1/list_explevels.txt \
--out eqtlbma_bf_quasipoisson_2000_85_112/eqtlbma_bf_quasipoisson_2000_85_112 \
--anchor TSS --cis 100000 --lik quasipoisson --analys join --outss --maf 0.05 \
--gridL /project/mstephens/gtex/analysis/june2014/types_v1/grid_phi2_oma2_general.txt.gz \
--gridS /project/mstephens/gtex/analysis/june2014/types_v1/grid_phi2_oma2_with-configs.txt.gz \
--bfs sin --error uvlr --fiterr 0.5 -v 1 \
--gcoord genes_2000_85_112.bed.gz --snp snps_2000_85_112.txt.gz
```

4.2 What Next?

After discussion with Matthew and William on June 22, 2015 we decide to pause this analysis (from experience of Pritchard's lab the gain from using GLM does not worth it's computational cost).

Chapter 5

Analyzing GTEx Using eqtlbma

5.1 Pre-processing

5.1.1 Extract input data from archive

BASH

```
for i in GTEx_Analysis_2015-01-12_eQTLInputFiles_covariates.tar.gz GTEx_Analysis_2015-01-12_eQTLInputFiles_snpMatrices.tar.gz GTEx_A\
nalysis_2015-01-12_eQTLInputFiles_geneLevelNormalizedExpressionMatrices.tar.gz GTEx_Analysis_2015-01-12_eQTLInputFiles_snpMatricesSu\
pplement.tar.gz; do
for j in `tar -tf $i`; do
echo $i $j
mkdir -p $(basename $i .tar.gz)
echo -e '#!/bin/bash\ntar -zxvf "$i" "$j" -O | gzip --best > "$(basename $i .tar.gz)"/"$j".gz\necho complete!' | sbatch -J E\
xtract."$i"."$j" -o $LogDir/Extract."$i"."$j".o%j
done
done
```

5.1.2 Gene TSS coordinates file

```
GetTSSCoords "$DataPrefix"eQTLInputFiles_genePositions.txt.gz $InputDir/tss_coords.bed.gz
```

5.1.3 SNP coordinates file

Unfortunately the release does not come with a list of SNPs (union) involved. Getting such a list is quite a heavy duty. Takes 10min to extract ID's in parallel and 50min to concatenate them into a unique list in bed format:

```
GetSNPCoords "$DataPrefix"eQTLInputFiles_snpMatrices $InputDir/snp_coords.bed.gz 1
GetSNPCoords "$DataPrefix"eQTLInputFiles_snpMatrices $InputDir/snp_coords.bed.gz 2
```

Note

In Sarah's analysis based on an earlier version of data, there are 55,993 genes and 6,856,776 SNPs provided in gene/snp lists. In the v6 release there are 56,318 genes and 10,297,646 SNPs listed; although the sumstats of v6 only has ~39,000 genes.

5.1.4 Input file lists

BASH

```
rm -rf $InputDir/list_geno.txt
DFolder="$DataPrefix"eQTLInputFiles_snpMatrices
for i in `ls $DFolder`; do
    j=`basename $i _Analysis.snps.txt.gz`
    echo -e "$j\t$DFolder/$i" >> $InputDir/list_geno.txt
done
#
rm -rf $InputDir/list_expr.txt
DFolder="$DataPrefix"eQTLInputFiles_geneLevelNormalizedExpressionMatrices
for i in `ls $DFolder`; do
    j=`basename $i _Analysis.expr.txt.gz`
    echo -e "$j\t$DFolder/$i" >> $InputDir/list_expr.txt
done
#
rm -rf $InputDir/list_covar.txt
DFolder="$DataPrefix"eQTLInputFiles_covariates
for i in `ls $DFolder`; do
    j=`basename $i _Analysis.covariates.txt.gz`
    echo -e "$j\t$DFolder/$i" >> $InputDir/list_covar.txt
done
```

5.2 The Configuration Model

5.2.1 Run eqtlbma_bf analysis in batches

analysis_admin eqtlbma_batch generates batches on the fly and perform eqtlbma_bf analysis in (embarrassing) parallel fashion on batches of genes. To use it,

```
python analysis_admin.py eqtlbma_batch -h
```

eqtlbma_batch

```
usage: analysis_admin.py eqtlbma_batch [-h] -g GENE_COORDS -s SNP_COORDS
                                         [-w N] [-n N_BATCHES] [-b BATCH_ID] -a
                                         ARGS_FILE [--seed N] -e EQTLBMA_PATH
                                         [--dry-run] [--clean]
```

optional arguments:

```
-h, --help            show this help message and exit
-g GENE_COORDS        Gene (TSS) coordinate file, in bed.gz format. (default:
                        None)
-s SNP_COORDS         SNP coordinate file, in bed.gz format. (default: None)
-w N                  Window size. (default: 100000)
-n N_BATCHES          Total number of batches. (default: 1000)
-b BATCH_ID           Execute the i-th batch. Program will quit if invalid ID is
                        provided. (default: 1)
-a ARGS_FILE          Path to file containing additional eqtlbma command
                        arguments. (default: None)
--seed N              If specified, a random number will be generated using (N +
                        batch_id) as seed, and the eqtlbma command will be appended
```

```

        a "--seed" argument with the number generated here.
        (default: None)
-e EQTLBMA_PATH Path to an eqtlbma.* executable. (default: None)
--dry-run       Only generate and save the batch data & commands without
                 performing analysis. (default: False)
--clean         Remove batch gene / snp coords file upon finishing the
                 analysis. (default: False)

```

Here is the command which breaks the data into 100 batches and perform analysis

100 batches analysis

```

nBatches=100
Model=normal
for i in `seq $nBatches`; do
echo '#!/bin/bash
source $HOME/GIT/gtex-eqtlsrc/cfg/GTE.bashrc
python $SrcDir/python/analysis_admin.py eqtlbma_batch \
-g $InputDir/tss_coords.bed.gz \
-s $InputDir/snp_coords.bed.gz \
-n "$nBatches" -b "$i" -w 100000 \
--seed 10086 -e "/software/bin/eqtlbma_bf \
-a $ConfDir/eqtlbma."$Model".txt' | \
sbatch -J eqtlbma_"$Model"_"$nBatches"_"$i" -e $LogDir/eqtlbma_"$Model"_"$nBatches"_"$i".e%j \
-o $LogDir/eqtlbma_"$Model"_"$nBatches"_"$i".o%j --mem-per-cpu=10000 --time=36:00:00
sleep 1
done

```

See my lab notebook for various failures and fixes running this command. Read on for procedure that succeeded.

5.2.2 Configuration model without permutation

Job submission

No permutation for configuration model

```

nBatches=100
Model=normal.perm0
for i in `seq $nBatches`; do
echo '#!/bin/bash
source $HOME/GIT/gtex-eqtlsrc/cfg/GTE.bashrc
python $SrcDir/python/analysis_admin.py eqtlbma_batch \
-g $InputDir/tss_coords.bed.gz \
-s $InputDir/snp_coords.bed.gz \
-n "$nBatches" -b "$i" -w 100000 \
--seed 10086 -e "/software/bin/eqtlbma_bf \
-a $ConfDir/eqtlbma."$Model".txt' | \
sbatch -J eqtlbma_"$Model"_"$nBatches"_"$i" -e $LogDir/eqtlbma_"$Model"_"$nBatches"_"$i".e%j \
-o $LogDir/eqtlbma_"$Model"_"$nBatches"_"$i".o%j --mem-per-cpu=10000 --time=36:00:00
sleep 1
done

```

Run completed under 24 hours, yielding to 45GB output file in various batches. For storage purpose it is best idea to first convert these files in different batches to HDF5 and merge them to one file & archive. Skipping this step for now as I want to move on to the HM and BMA steps. These output are archived as is, though:

```
tar -cvf eqtlbma_bf.normal.perm0.tar *
```

5.2.3 List of summary statistics

The following script produces sumstats lists for all results previously computed:

make sumstats lists

```
SumstatsDir=$ArchiveDir/eqlbma_bf/July2015
for i in `find $SumstatsDir -maxdepth 1 -name "eqlbma_bf_normal_*" -type d`; do
  for j in `ls $i/*_sumstats_*.txt.gz`; do
    echo `echo $j | sed 's/\(.*sumstats_\)\(.*\).txt.gz/\2/g` $j >> $i.ss.list
  done
done
```

Future analysis can be based on summary statistics instead.

5.3 The Type Model

The raw BF data are fed to the type-model version of the hierarchical model EM algorithm to estimate hyperparameters. Data is 28GB compressed text which may exceed 256GB when all are loaded into RAM (which is what eqlbma_hm does!).

See my labnotes of July, 2015 to August, 2015 for failures and fixes.